

Tutorial Human Genomeme Annotation

1. Introduction

1.1 What is gene annotation?

Over the past years, we have learnt that there are a number of chromosomes and genes in our genome. Counting the number of chromosomes is fairly easy but students might find difficult to say how many genes we have in our genome. If you can get an answer for this, could you tell how many genes encode protein and how many do not?

To answer this question, we need to access the database for gene annotation. Gene annotation is the process of making nucleotide sequence meaningful - where genes are located? whether it is protein-coding or noncoding. If you would like to get an overview of gene annotation, please find this link.

One of well-known collaborative efforts in gene annotation is the GENCODE consortium. It is a part of the Encyclopedia of DNA Elements (The ENCODE project consortium) and aims to identify all gene features in the human genome using a combination of computational analysis, manual annotation, and experimental validation (Harrow et al. 2012). You might find another database for gene annotation, like RefSeq, CCDS, and need to understand differences between them.

In this tutorial, we will access to gene annotation from the GENCODE consortium and explore genes and functional elements in our genome.

1.2 Aims

What we will do with this dataset: - Be familiar with gene annotation modality. - Tidy data and create a table for your analysis. - Apply tidyverse functions for data munging.

Please note that there is better solution for getting gene annotation in R if you use a biomart. Our tutorial is only designed to have a practice on tidyverse exercise.

2. Explore your data

2.1 Unboxing your dataset

This tutorial will use a gene annotation file from the GENCODE. You will need to download the file from the GENCODE. If you are using terminal, please download file using wget: Once you download the file, you can print out the first few lines using the following bash command (we will learn UNIX commands later):

The file is the GFT file format, which you will find most commonly in gene annotation. Please read the file format thoroughly in the link above.

For the tutorial, we need to load two packages. If the package is not installed in your system, please install it.

- tidyverse, a package you have learnt from the chapter 5.

- readr, a package provides a fast and friendly way to read. Since the file gencode.v31.basic.annotation.gtf.gz is pretty large, you will need some function to load data quickly into your workspace. readr is a part of tidyverse, so you can just load tidyverse to use readr functions.

Let's load the GTF file into your workspace. We will use read_delim function from the readr package. This is much faster loading than read.delim or read.csv from R base. However, please keep in mind that some parameters and output class for read_delim are slightly different from them.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.5      v dplyr  1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

d = read_delim("gencode.v31.basic.annotation.gtf.gz",
               delim = '\t', skip = 5, progress = F,
               col_names = F)

## Rows: 1756502 Columns: 9

## -- Column specification -----
## Delimiter: "\t"
## chr (7): X1, X2, X3, X6, X7, X8, X9
## dbl (2): X4, X5

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Can you find out what the parameters mean? Few things to note are: - The GTF file contains the first few lines for comments (#). In general, the file contains description, provider, date, format. - The GTF file does not have column names so you will need to assign FALSE for col_names.

This is sort of canonical way to load your dataset into R. However, we are using a GTF format, which is specific to gene annotation so we can use a package to specifically handle a GTF file.

Here I introduce the package rtracklayer. Let's install the package first.

```
if (!requireNamespace("BiocManager", quietly = TRUE))
  install.packages("BiocManager")

BiocManager::install("rtracklayer")
```

```
## Bioconductor version 3.13 (BiocManager 1.30.16), R 4.1.0 (2021-05-18)
```

```
## Warning: package(s) not installed when version(s) same as current; use `force = TRUE` to
## re-install: 'rtracklayer'
```

Then, now you can read the GTF file using this package. Then, you can check the class of the object d.

```
d = rtracklayer::import('gencode.v31.basic.annotation.gtf.gz')
class(d)
```

```
## [1] "GRanges"
## attr(,"package")
## [1] "GenomicRanges"
```

You will find out that this is GRanges class. This is from the package Genomic Range, specifically dealing with genomic datasets but we are not heading into this in this tutorial. So please find this information if you are serious on this.

We are converting d into a data frame as following:

```
d = d %>% as.data.frame()
```

Let's overview few lines from the data frame, and explore what you get in this object.

```
head(d)
```

```
##      seqnames start   end width strand source      type score phase
## 1      chr1 11869 14409   2541      + HAVANA      gene    NA     NA
## 2      chr1 11869 14409   2541      + HAVANA transcript NA     NA
## 3      chr1 11869 12227    359      + HAVANA      exon    NA     NA
## 4      chr1 12613 12721    109      + HAVANA      exon    NA     NA
## 5      chr1 13221 14409   1189      + HAVANA      exon    NA     NA
## 6      chr1 12010 13670   1661      + HAVANA transcript NA     NA
##      gene_id      gene_type gene_name level
## 1 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
## 2 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
## 3 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
## 4 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
## 5 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
## 6 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1      2
##      hgnc_id      havana_gene      transcript_id
## 1 HGNC:37102 OTTHUMG00000000961.2      <NA>
## 2 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 3 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 4 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 5 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 6 HGNC:37102 OTTHUMG00000000961.2 ENST00000450305.2
##      transcript_type transcript_name transcript_support_level
## 1      <NA>      <NA>      <NA>
## 2      lncRNA      DDX11L1-202      1
## 3      lncRNA      DDX11L1-202      1
## 4      lncRNA      DDX11L1-202      1
## 5      lncRNA      DDX11L1-202      1
## 6 transcribed_unprocessed_pseudogene DDX11L1-201      NA
##      tag      havana_transcript exon_number      exon_id      ont
```

```
## 1 <NA> <NA> <NA> <NA> <NA>
## 2 basic OTTHUMT00000362751.1 <NA> <NA> <NA>
## 3 basic OTTHUMT00000362751.1 1 ENSE00002234944.1 <NA>
## 4 basic OTTHUMT00000362751.1 2 ENSE00003582793.1 <NA>
## 5 basic OTTHUMT00000362751.1 3 ENSE00002312635.1 <NA>
## 6 basic OTTHUMT00000002844.2 <NA> <NA> PGO:0000019
## protein_id ccidsid
## 1 <NA> <NA>
## 2 <NA> <NA>
## 3 <NA> <NA>
## 4 <NA> <NA>
## 5 <NA> <NA>
## 6 <NA> <NA>
```

One thing you can find is that there is no columns in the data frame. Let's match which information is provided in columns. You can find the instruction page in the website ([link](#)).

Based on this, you can assign a name for 9 columns. One thing to remember is you should not use space for the column name. Spacing in the column name is actually working but not a good habit for your code. So please replace a space with underscore in the column name.

```
# Assign column names according to the GENCODE instruction.
cols = c('chrom', 'source', 'feature_type', 'start', 'end', 'score', 'strand', 'phase', 'info')
```

Now you can set up the column names into the col_names parameter, and load the file into a data frame.

```
d = read_delim('gencode.v31.basic.annotation.gtf.gz',
               delim='\t', skip = 5,
               progress = F,
               col_names = cols)
```

```
## Rows: 1756502 Columns: 9
```

```
## -- Column specification -----
## Delimiter: "\t"
## chr (7): chrom, source, feature_type, score, strand, phase, info
## dbl (2): start, end

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

You can find the column names are now all set.

```
head(d)

## # A tibble: 6 x 9
##   chrom source feature_type start   end score strand phase info
##   <chr> <chr>   <chr>      <dbl> <dbl> <chr> <chr>   <chr> <chr>
## 1 chr1  HAVANA  gene        11869 14409 .    +    .    "gene_id \"ENSG000000~
## 2 chr1  HAVANA  transcript   11869 14409 .    +    .    "gene_id \"ENSG000000~
## 3 chr1  HAVANA  exon        11869 12227 .    +    .    "gene_id \"ENSG000000~
## 4 chr1  HAVANA  exon        12613 12721 .    +    .    "gene_id \"ENSG000000~
## 5 chr1  HAVANA  exon        13221 14409 .    +    .    "gene_id \"ENSG000000~
## 6 chr1  HAVANA  transcript   12010 13670 .    +    .    "gene_id \"ENSG000000~
```

When you loaded the file, you see the message about the data class. You might want to overview this data.

```
summary(d)
```

```
##      chrom      source      feature_type      start
## Length:1756502 Length:1756502 Length:1756502 Min. :      577
## Class :character Class :character Class :character 1st Qu.: 32101517
## Mode :character Mode :character Mode :character Median : 61732754
##                                     Mean : 75288563
##                                     3rd Qu.:111760181
##                                     Max. :248936581
##      end      score      strand      phase
## Min. :      647 Length:1756502 Length:1756502 Length:1756502
## 1st Qu.: 32107331 Class :character Class :character Class :character
## Median : 61738373 Mode :character Mode :character Mode :character
## Mean : 75292632
## 3rd Qu.:111763007
## Max. :248937043
##      info
## Length:1756502
## Class :character
## Mode :character
##
##
##
```

2.2 How many feature types in the GENCODE dataset?

As instructed in the GENCODE website, the GENCODE dataset provides a range of annotations for the feature type. You can check feature types using _____ function.

```
d %>% group_by(feature_type) %>% count(feature_type)
```

```
## # A tibble: 8 x 2
## # Groups:   feature_type [8]
##   feature_type      n
##   <chr>      <int>
## 1 CDS        567862
## 2 exon       744835
## 3 gene        60603
## 4 Selenocysteine 96
## 5 start_codon   57886
## 6 stop_codon    57775
## 7 transcript    108243
## 8 UTR          159202
```

```
table(d$feature_type)
```

```
##
##      CDS      exon      gene Selenocysteine      start_codon
##      567862      744835      60603      96      57886
##      stop_codon      transcript      UTR
##      57775      108243      159202
```

How many feature types provided in the GENCODE? And how many items stored for each feature type? Please write down the number of feature types from the dataset. Also, if you are not familiar with these types, it would be good to put one or two sentences that can describe each type).

```
# 8 feature types
```

2.3 How many genes we have?

Let's count the number of genes in our genome. Since we know that the column `feature_type` contains rows with `gene`, which contains obviously annotations for genes. We might want to subset those rows from the data frame.

```
d1 <- filter(d, feature_type == "gene")
nrow(d1)
```

```
## [1] 60603
```

2.4 Ensembl, Havana and CCDS.

Gene annotation for the human genome is provided by multiple organizations with different gene annotation methods and strategy. This means that information can be varying by resources, and users need to understand heterogeneity inherent in annotation databases.

The GENCODE project utilizes two sources of gene annotation. 1. Havana: Manual gene annotation (detailed strategy in here) 2. Ensembl: Automatic gene annotation (detailed strategy in here)

It provides the combination of Ensembl/HAVANA gene set as the default gene annotation for the human genome. In addition, they also guarantee that all transcripts from the Consensus Coding Sequence (CCDS) set are present in the GENCODE gene set. The CCDS project is a collaborative effort to identify a core set of protein coding regions that are consistently annotated and of high quality. Initial results from the Consensus CDS (CCDS) project are now available through the appropriate Ensembl gene pages and from the CCDS project page at NCBI. The CCDS set is built by consensus among Ensembl, the National Center for Biotechnology Information (NCBI), and the HUGO Gene Nomenclature Committee (HGNC) for human (link).

Right. Then now we count how many genes annotated with HAVANA and ENSEMBL.

```
d %>% group_by(source) %>% count(source)
```

```
## # A tibble: 2 x 2
## # Groups:   source [2]
##   source      n
##   <chr>    <int>
## 1 ENSEMBL 245185
## 2 HAVANA 1511317
```

2.5 do.call

Since the last column `info` contains a long string for multiple annotations, we will need to split it to extract each annotation. For example, the first line for transcript annotation looks like this:

```
# chr1 HAVANA transcript 11869 14409 . + . gene_id "ENSG00000223972.5"; transcr
```

If you would like to split transcript_support_level and create a new column, you can use strsplit function.

```
a = 'chr1 HAVANA transcript 11869 14409 . + . gene_id "ENSG00000223972.5"; transcr
strsplit(a, 'transcript_support_level\\s+')

## [[1]]
## [1] "chr1 HAVANA transcript 11869 14409 . + . gene_id \"ENSG00000223972.5\"; transcr
## [2] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc
```

After split the string, you can select the second item in the list ([[1]][2]).

```
strsplit(a, 'transcript_support_level\\s+')[[1]][2]

## [1] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc
```

You can find the 1 in the first position, which you will need to split again.

```
b = strsplit(a, 'transcript_support_level\\s+')[[1]][2]
strsplit(b, '\\\\')

## [[1]]
## [1] "1" " "; hgnc_id " " "HGNC:37102"
## [4] " "; tag " " "basic" " "; havana_gene "
## [7] "OTTHUMG00000000961.2" " "; havana_transcript " "OTTHUMT00000362751.1"
## [10] " ;"
```

From this, you will get the first item in the list ([[1]][1]).

Now you would like to apply strsplit function across vectors. For this, do.call function can be easily implemented to strsplit over the vectors from one column. Let's try this.

```
head(do.call(rbind.data.frame, strsplit(a, 'transcript_support_level\\s+'))[[2]])

## [1] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc
```

Now you can write two lines of codes to process two steps we discussed above.

```
# First filter transcripts and create a data frame.
d2 <- d %>% filter(feature_type == 'transcript')

# Now apply the functions.
d2$transcript_support_level <- as.character(do.call(rbind.data.frame,
                                                    strsplit(d2$info, 'transcript_support_level\\s+'))
d2$transcript_support_level <- as.character(do.call(rbind.data.frame,
                                                    strsplit(d2$transcript_support_level, '\\\\')[[1]]
```

Now you can check the strsplit works.

```
head(d2$transcript_support_level)
```

```
## [1] "1" "NA" "NA" "NA" "5" "5"
```

You can use the same method to extract other annotations, like `gene_id`, `gene_name` etc.

3. Exercises

Here I list the questions for your activity. Please note that it is an exercise for tidyverse functions, which you will need to use in your code. In addition, you will need to write an one-line code for each question using pipe `%>%`.

For questions, you should read some information thoroughly, including:

- Gene biotype.
- 0 or 1 based annotation in GTF, BED format
- Why some features have 1 bp length?
- What is the meaning of zero-length exons in GENCODE? Also fun to have a review for microexons
- Transcript support level (TSL)

```
d_ <- rtracklayer::import('gencode.v31.basic.annotation.gtf.gz')
d_ <- as.data.frame(d_)
d_ <- d_ %>% rename(chromosome = seqnames)
```

3.1 Annotation of transcripts in our genome

Q1. Computes the number of transcripts per gene. What is the mean number of transcripts per gene? What is the quantile (25%, 50%, 75%) for these numbers? Which gene has the greatest number of transcript?

```
# <Mean number of transcripts per gene>
d_2 <- d_ %>%
  filter(type == "transcript")
tperg <- d_2 %>%
  count(gene_id)
tperg %>%
  summarize(mean = mean(n))
```

```
##      mean
## 1 1.7861
```

```
# <Quantile>
p <- seq(0.25, 0.75, 0.25)
quantile(tperg$n, p)
```

```
## 25% 50% 75%
##   1   1   2
```



```
# <Max>
max(tperg$n)
```

```
## [1] 87
```

Q2. Compute the number of transcripts per gene among gene biotypes. For example, compare the number of transcript per gene between protein-coding genes, long noncoding genes, pseudogenes.

```
# <Mean number of transcripts per gene among gene biotypes>
# Count the number of transcripts per gene among gene biotypes
tperg_type <- d_2 %>% group_by(gene_type) %>% summarize(gene_n = n_distinct(gene_id), transcripts_n=n()
tperg_type
```

```
## # A tibble: 40 x 4
##   gene_type      gene_n transcripts_n t_per_g
##   <chr>          <int>         <int>   <dbl>
## 1 IG_C_gene       14           14     1
## 2 IG_C_pseudogene  9            9     1
## 3 IG_D_gene      37           37     1
## 4 IG_J_gene      18           18     1
## 5 IG_J_pseudogene  3            3     1
## 6 IG_pseudogene   1            1     1
## 7 IG_V_gene     144          144     1
## 8 IG_V_pseudogene 188          188     1
## 9 lncRNA       16840         24993   1.48
## 10 miRNA        1881         1881     1
## # ... with 30 more rows
```

```
# Protein-coding genes, long noncoding genes, pseudogenes
d_2 <- d_2 %>%
  mutate(group = case_when(
    gene_type %in% c("IG_C_gene", "IG_D_gene", "IG_J_gene", "IG_LV_gene", "IG_V_gene", "TR_C_gene", "TR_J_gene") ~ "Protein-coding genes",
    gene_type %in% c("IG_pseudogene", "IG_C_pseudogene", "IG_J_pseudogene", "IG_V_pseudogene", "TR_V_pseudogene") ~ "pseudogenes",
    gene_type == "lncRNA" ~ "long noncoding gene",
    TRUE ~ "Others"
  ))
```

```
d_2 %>% group_by(group) %>% summarize(gene_n = n_distinct(gene_id), transcripts_n=n(), t_per_g = transcripts_n/gene_n)
```

```
## # A tibble: 4 x 4
##   group          gene_n transcripts_n t_per_g
##   <chr>          <int>         <int>   <dbl>
## 1 long noncoding gene 16840         24993   1.48
## 2 Others            8140         8140     1
## 3 Protein-coding genes 20383        58254   2.86
## 4 pseudogenes       15240        16856   1.11
```

Q3. Final task is to compute the number of transcripts per gene per chromosome.

```
# Count the number of transcripts per gene per chromosome
tperg_chrm <- d_2 %>% group_by(chromosome) %>% summarize(transcripts_n=n(), gene_n = n_distinct(gene_id))
tperg_chrm
```

```
## # A tibble: 25 x 4
##   chromosome transcripts_n gene_n t_per_g
##   <fct>          <int>  <int>   <dbl>
## 1 chr1           9827   5471    1.80
## 2 chr2           7432   4196    1.77
## 3 chr3           6157   3185    1.93
## 4 chr4           4662   2651    1.76
## 5 chr5           5203   2983    1.74
## 6 chr6           5455   3059    1.78
## 7 chr7           5292   3014    1.76
## 8 chr8           4350   2482    1.75
## 9 chr9           3949   2327    1.70
## 10 chr10          4157   2332    1.78
## # ... with 15 more rows
```

3.2 Gene length in the GENCODE

Q1. What is the average length of human genes?

```
d_1 <- d_1 %>%
  filter(type == "gene")
mean(d_1$width)
```

```
## [1] 32629.02
```

Q2. Is the distribution of gene length differed by autosomal and sex chromosomes? Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for each group.

```
# 1) Add "chrom_group" column
d_1 <- d_1 %>%
  mutate(chrom_group = case_when(
    chromosome == "chrM" ~ "Mitrochondrial",
    chromosome %in% c("chrX", "chrY") ~ "Sex",
    TRUE ~ "Autosomal" ))

# 2) Quantiles of gene length for each group
Sex_length <- d_1 %>%
  filter(chrom_group == "Sex") %>%
  .$width

Autosome_length <- d_1 %>%
  filter(chrom_group == "Autosomal") %>%
  .$width

quantile(Sex_length,)
```

```
##      0%      25%      50%      75%     100%
##      48      473     1912     13502  2241765
```

```
quantile(Autosome_length,)
```

```
##      0%      25%      50%      75%     100%
##      8      565     3779     25813  2473537
```

Q3. Is the distribution of gene length differed by gene biotype? Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for each group.

```
d_1 %>%
  group_by(gene_type) %>%
  summarize(quant0 = quantile(width, probs = 0),
            quant25 = quantile(width, probs = 0.25),
            quant50 = quantile(width, probs = 0.5),
            quant75 = quantile(width, probs = 0.75),
            quant100 = quantile(width, probs = 1))
```

```
## # A tibble: 40 x 6
##   gene_type      quant0 quant25 quant50 quant75 quant100
##   <chr>          <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 IG_C_gene      441     477.    4590.   5479.    8914
## 2 IG_C_pseudogene 248     313     317     734     5211
## 3 IG_D_gene       11      17      20      31       37
## 4 IG_J_gene       33     38.2     49     67.8    176
## 5 IG_J_pseudogene  50     52.5     55     57.5     60
## 6 IG_pseudogene   306    306     306    306     306
## 7 IG_V_gene      294    496.     522    572.   176628
## 8 IG_V_pseudogene  28     271     416.    458.     792
## 9 lncRNA          68   1874.    6272.  24774.  1375317
## 10 miRNA          41      70      80      91      180
## # ... with 30 more rows
```

3.3 Transcript support levels(TSL)

The GENCODE TSL provides a consistent method of evaluating the level of support that a GENCODE transcript annotation is actually expressed in humans.

Q1. With transcript, how many transcripts are categorized for each TSL?

```
d_2 %>% group_by(transcript_support_level) %>% count()
```

```
## # A tibble: 7 x 2
## # Groups:   transcript_support_level [7]
##   transcript_support_level     n
##   <chr>                   <int>
## 1 1                       31801
## 2 2                       13372
## 3 3                        7228
## 4 4                        2245
## 5 5                       13674
## 6 NA                      27843
## 7 <NA>                    12080
```

Q2. From the first question, please count the number of transcript for each TSL by gene biotype.

```
d_2 %>% group_by(gene_type, transcript_support_level) %>% count()
```

```
## # A tibble: 91 x 3
## # Groups:   gene_type, transcript_support_level [91]
##   gene_type      transcript_support_level     n
##   <chr>         <chr>                   <int>
## 1 IG_C_gene      1                       1
## 2 IG_C_gene      5                       1
## 3 IG_C_gene      NA                      7
## 4 IG_C_gene      <NA>                   5
## 5 IG_C_pseudogene NA                      9
## 6 IG_D_gene      NA                     37
## 7 IG_J_gene      NA                     18
## 8 IG_J_pseudogene NA                      3
## 9 IG_pseudogene   NA                      1
##10 IG_V_gene      5                      3
## # ... with 81 more rows
```

Q3. From the first question, please count the number of transcript for each TSL by source.

```
d_2 %>% group_by(source, transcript_support_level) %>% count()
```

```
## # A tibble: 14 x 3
## # Groups:   source, transcript_support_level [14]
##   source transcript_support_level     n
##   <fct>   <chr>                   <int>
## 1 HAVANA  1                29434
## 2 HAVANA  2                12052
## 3 HAVANA  3                 6964
## 4 HAVANA  4                 2116
## 5 HAVANA  5                10157
## 6 HAVANA  NA                19962
## 7 HAVANA  <NA>             11901
## 8 ENSEMBL 1                 2367
## 9 ENSEMBL 2                 1320
##10 ENSEMBL 3                  264
##11 ENSEMBL 4                  129
##12 ENSEMBL 5                 3517
##13 ENSEMBL NA                7881
##14 ENSEMBL <NA>              179
```

3.4 CCDS in the GENCODE

Q1. With gene, please create a data frame with the columns - gene_id, gene_name, hgnc_id, gene_type, chromosome, start, end, and strand. Then, please create new columns for presence of hgnc and ccds. For example, you can put 1 in the column isHgnc, if hgnc annotation is available, or 0 if not. Then, you can put 1 in the column isCCDS, if ccds annotation is available, or 0 if not.

```

# Create a new data frame
d_gene <- d_1 %>%
  select(gene_id, gene_name, hgnc_id, gene_type, chromosome, start, end, strand)

# column isHgnc
d_gene <- d_gene %>%
  mutate(isHgnc = case_when(
    is.na(hgnc_id) ~ 0,
    TRUE ~ 1
  ))

# column isCCDS
CCDS_gene <- d_ %>%
  filter(!is.na(ccdsid)) %>%
  count(gene_id) %>%
  .$gene_id

d_gene <- d_gene %>%
  mutate(isCCDS = case_when(
    gene_id %in% CCDS_gene ~ 1,
    TRUE ~ 0
  ))

head(d_gene)

```

```

##           gene_id  gene_name  hgnc_id  gene_type
## 1 ENSG00000223972.5    DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
## 2 ENSG00000227232.5    WASH7P HGNC:38034      unprocessed_pseudogene
## 3 ENSG00000278267.1  MIR6859-1 HGNC:50039                miRNA
## 4 ENSG00000243485.5 MIR1302-2HG HGNC:52482                lncRNA
## 5 ENSG00000284332.1  MIR1302-2 HGNC:35294                miRNA
## 6 ENSG00000237613.2    FAM138A HGNC:32334                lncRNA
##   chromosome start  end strand isHgnc isCCDS
## 1      chr1 11869 14409      +      1      0
## 2      chr1 14404 29570      -      1      0
## 3      chr1 17369 17436      -      1      0
## 4      chr1 29554 31109      +      1      0
## 5      chr1 30366 30503      +      1      0
## 6      chr1 34554 36081      -      1      0

```

Q2. Please count the number of hgnc by gene biotypes.

```

d_gene %>%
  group_by(gene_type) %>%
  summarize(hgnc_id_n = n_distinct(hgnc_id))

```

```

## # A tibble: 40 x 2
##   gene_type  hgnc_id_n
##   <chr>      <int>
## 1 IG_C_gene      14
## 2 IG_C_pseudogene 9
## 3 IG_D_gene     37

```

```
## 4 IG_J_gene 18
## 5 IG_J_pseudogene 3
## 6 IG_pseudogene 1
## 7 IG_V_gene 143
## 8 IG_V_pseudogene 186
## 9 lncRNA 3962
## 10 miRNA 1855
## # ... with 30 more rows
```

Q3. Please count the number of hgnc by level. Please note that level in this question is not TSL. Please find information in this link: 1 (verified loci), 2 (manually annotated loci), 3 (automatically annotated loci).

```
d_gene %>%
  mutate(level = d_1$level) %>%
  group_by(level) %>%
  summarize(hgnc_id_n = n_distinct(hgnc_id))
```

```
## # A tibble: 3 x 2
##   level hgnc_id_n
##   <chr>   <int>
## 1 1      11342
## 2 2      20475
## 3 3       5725
```

3.5 Transcripts in the GENCODE

Q1. Which gene has the largest number of transcripts?

```
d_ %>%
  group_by(gene_id, gene_name) %>%
  count() %>%
  as.data.frame() %>%
  filter(n == max(n))
```

```
##           gene_id gene_name    n
## 1 ENSG00000155657.26      TTN 3876
```

Q2. Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for protein coding genes and long noncoding genes.

```
d_1 %>%
  filter(gene_type %in% c("lncRNA", "protein_coding")) %>%
  group_by(gene_type) %>%
  summarize(quant0 = quantile(width, probs = 0),
            quant25 = quantile(width, probs = 0.25),
            quant50 = quantile(width, probs = 0.5),
            quant75 = quantile(width, probs = 0.75),
            quant100 = quantile(width, probs = 1))
```

```
## # A tibble: 2 x 6
##   gene_type quant0 quant25 quant50 quant75 quant100
```

```
##   <chr>           <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 lncRNA           68    1874.    6272.   24774.  1375317
## 2 protein_coding   117    9632.   27212   70809   2473537
```

Q3. Please count the number of transcripts by chromosomes.

```
d_2 %>%
  group_by(chromosome) %>%
  count()
```

```
## # A tibble: 25 x 2
## # Groups:   chromosome [25]
##   chromosome     n
##   <fct>       <int>
## 1 chr1         9827
## 2 chr2         7432
## 3 chr3         6157
## 4 chr4         4662
## 5 chr5         5203
## 6 chr6         5455
## 7 chr7         5292
## 8 chr8         4350
## 9 chr9         3949
## 10 chr10        4157
## # ... with 15 more rows
```

3.6 Autosomal vs. Sex chromosomes

Q1. Please calculate the number of genes per chromosome.

```
d_1 %>%
  group_by(chromosome) %>%
  count()
```

```
## # A tibble: 25 x 2
## # Groups:   chromosome [25]
##   chromosome     n
##   <fct>       <int>
## 1 chr1         5471
## 2 chr2         4196
## 3 chr3         3185
## 4 chr4         2651
## 5 chr5         2983
## 6 chr6         3059
## 7 chr7         3014
## 8 chr8         2482
## 9 chr9         2327
## 10 chr10        2332
## # ... with 15 more rows
```

Q2. Please compare the number of genes between autosomal and sex chromosome (Mean, Median).

chrom_group in d_1 already has been defined in previous questions

```
d_1 %>%
  filter(chrom_group %in% c("Autosomal", "Sex")) %>%
  group_by(chrom_group, chromosome) %>%
  count() %>%
  group_by(chrom_group) %>%
  summarize(sum = sum(n), mean = mean(n), median = median(n))
```

```
## # A tibble: 2 x 4
##   chrom_group    sum  mean median
##   <chr>         <int> <dbl>  <dbl>
## 1 Autosomal   57577 2617.  2604.
## 2 Sex         2989 1494.  1494.
```

Q3. Please divide the genes into groups 'protein coding' and 'long noncoding', and then compare the number of genes in each chromosomes within groups.

```
d_1 %>%
  filter(gene_type %in% c("protein_coding", "lncRNA")) %>%
  group_by(gene_type, chromosome) %>%
  count()
```

```
## # A tibble: 49 x 3
## # Groups:   gene_type, chromosome [49]
##   gene_type chromosome      n
##   <chr>      <fct>      <int>
## 1 lncRNA    chr1         1416
## 2 lncRNA    chr2         1241
## 3 lncRNA    chr3          861
## 4 lncRNA    chr4          790
## 5 lncRNA    chr5          950
## 6 lncRNA    chr6          826
## 7 lncRNA    chr7          720
## 8 lncRNA    chr8          831
## 9 lncRNA    chr9          555
## 10 lncRNA   chr10         695
## # ... with 39 more rows
```